

Introduction to Deep Learning

1. Logistics, Software, Linear Algebra

STAT 157, Spring 2019, UC Berkeley

Mu Li and Alex Smola

courses.d2l.ai/berkeley-stat-157

Logistics



Goals

- **Introduction** to Deep Learning
(basic MLP, optimization, convolutions, sequences)
- **Theory**
 - Capacity control (weight decay, dropout, batch norm)
 - Optimization, models, overfitting, objective functions
- **Practice**
 - Write code in Python / MxNet Gluon / Jupyter
 - Solve realistic problems
- **Project**
 - Ability to start original research in Deep Learning **in a team**

Getting there

- **Course**
 - Tuesday and Thursday 3:30-5:00pm in LeConte 3
 - Slides and notebooks online on website
 - YouTube channel with lectures at goo.gl/zcBh2U
 - Github repository at [d2l-ai/berkeley-stat-157](https://github.com/d2l-ai/berkeley-stat-157)
- **Dive into Deep Learning**
 - Jupyter Notebooks
 - Github repository at [d2l-ai/d2l-en](https://github.com/d2l-ai/d2l-en)

Contacts

- **Lecturers**

- Mu Li and Alex Smola
Office hours Thursdays 1:00-2:30pm in Evans Hall 337
 - Email berkeley-stat-157@googlegroups.com

- **TA**

- Ryan Theisen
Office hours Wednesdays, likely 2:00-4:00pm in Evans Hall
 - Email theisen@berkeley.edu
- Course discussion site discuss.mxnet.io/c/courses

Homework

- **10 assignments**
 - Due 1 week after posted
1/29, 2/5, 2/12, 2/19, 2/26, 3/12, 4/2, 4/9, 4/16, 4/30
 - Solutions and feedback returned 1 week after submitted.
 - Solutions based on students' submissions (we will acknowledge the writer)
- **Best 9 out of 10** homeworks count
- You can be **late by 2 days once** for any reason
- 50/50 mix of programming and theory

Homework

- **Submit homework via GitHub**
 - Submit the homework by 4pm on the Tuesday it's due
 - TA performs pull request after deadline
 - Submit as PDF (theory) and Jupyter notebooks (code)
 - TA commits annotated feedback via Git - add **rythei**
- **Logistics**
 - Github account & repository (email to course)
 - Permission for TA to read/write the repository

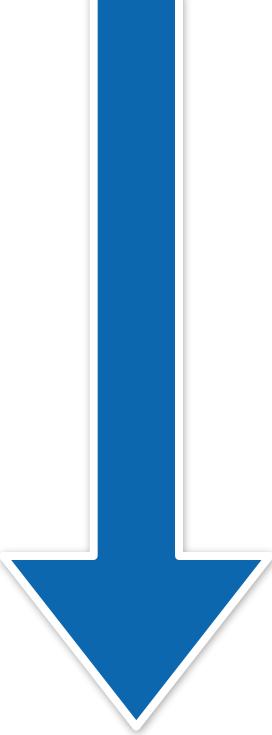
Exam

- **March 19, 2019**
- Midterm exam only
 - Open book but not open computer/phone/tablet
 - Problems will be similar to homework
- **Why**
 - At least one test that has to be done in person
(you *could* fake the rest)
 - No final exam due to projects

Project

- **Original work in machine learning**
 - Existing tools applied to novel problem
 - Novel tools
 - Ask lecturers & TA for advice if you're stuck
- Research 'with training wheels' **simulates academic process**
 - Research in a team (**3-5 students**)
 - Deliverables with schedule / deadlines
 - End result is a paper/report/presentation (NIPS template)

Project

- 
- 2/5 **Register team** (names, working title)
 - 3/5 **Project proposal** (1-2 page, 5 min talk)
 - 4/22-25 (or earlier) **Talk to TA** to discuss
 - 5/7-9 **Final presentation & report**
(6-20 pages report, 6-20 slides talk)
 - Talk to TA and lecturers (it's our job to help)
 - Start **early** (last minute projects fail often)
 - No, you cannot do it alone. **This is teamwork.**

Have fun. Learn stuff.

Deep Learning

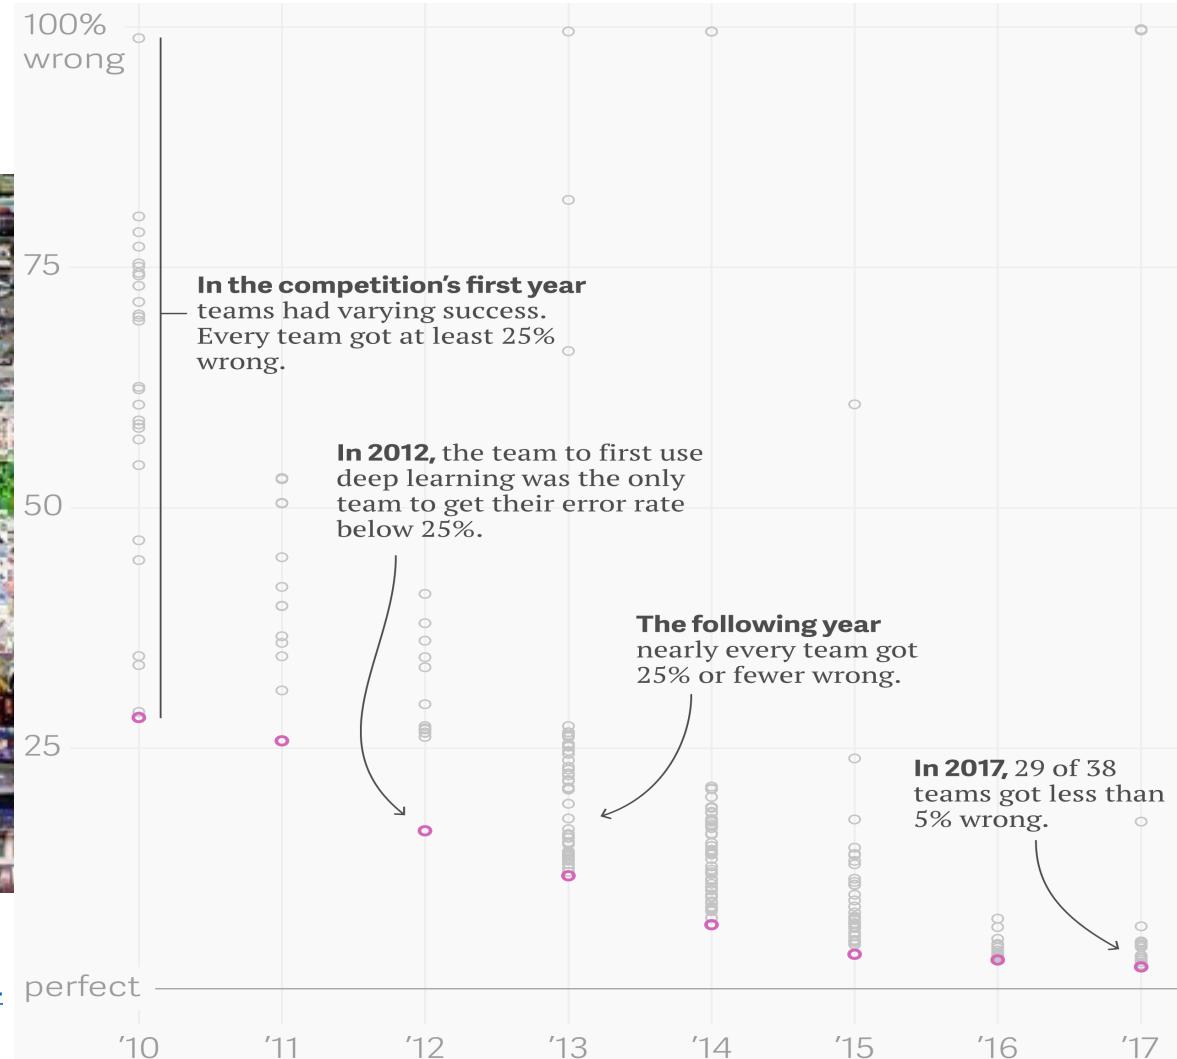
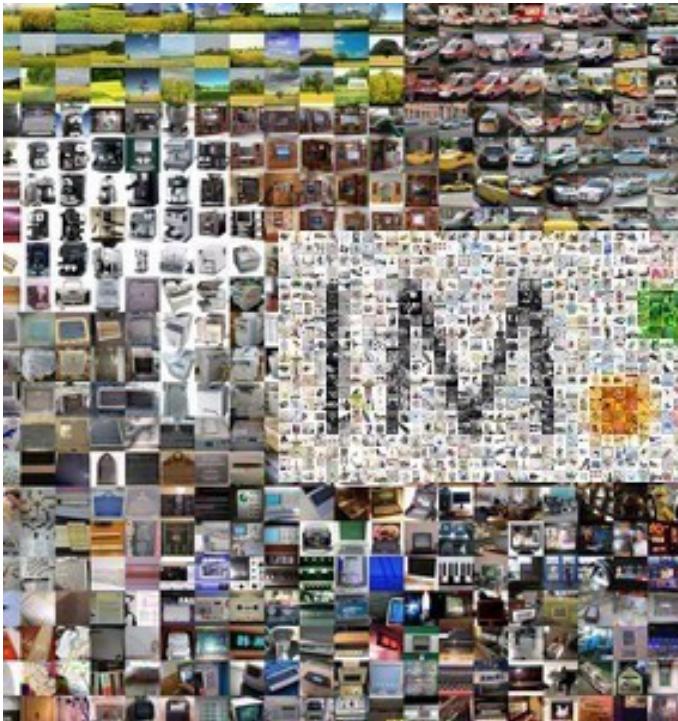


Classify Images



<http://www.image-net.org/>

Classify Images



Yanofsky, Quartz

<https://qz.com/1034972/the-data-that-changed-the-direction-of-ai-research-and-possibly-the-world/>

Detect and Segment Objects



Style transfer

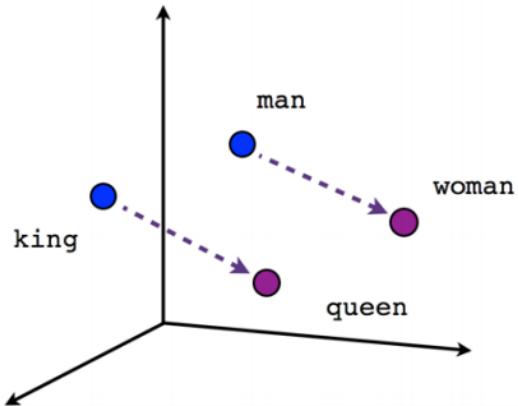


<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>

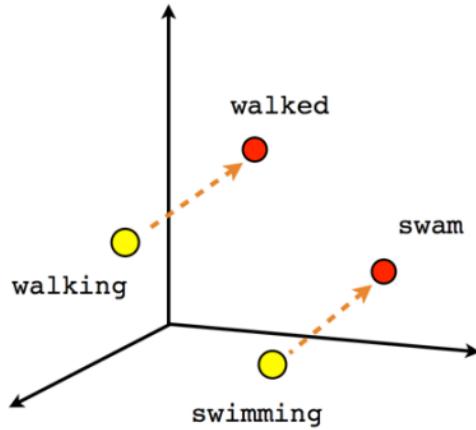
Synthesize Faces



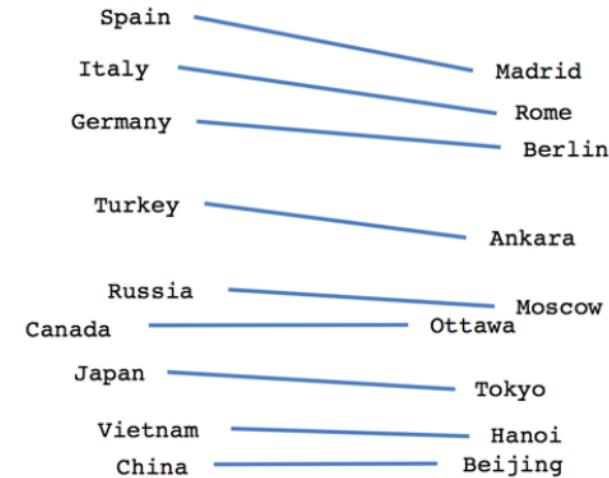
Analogies



Male-Female



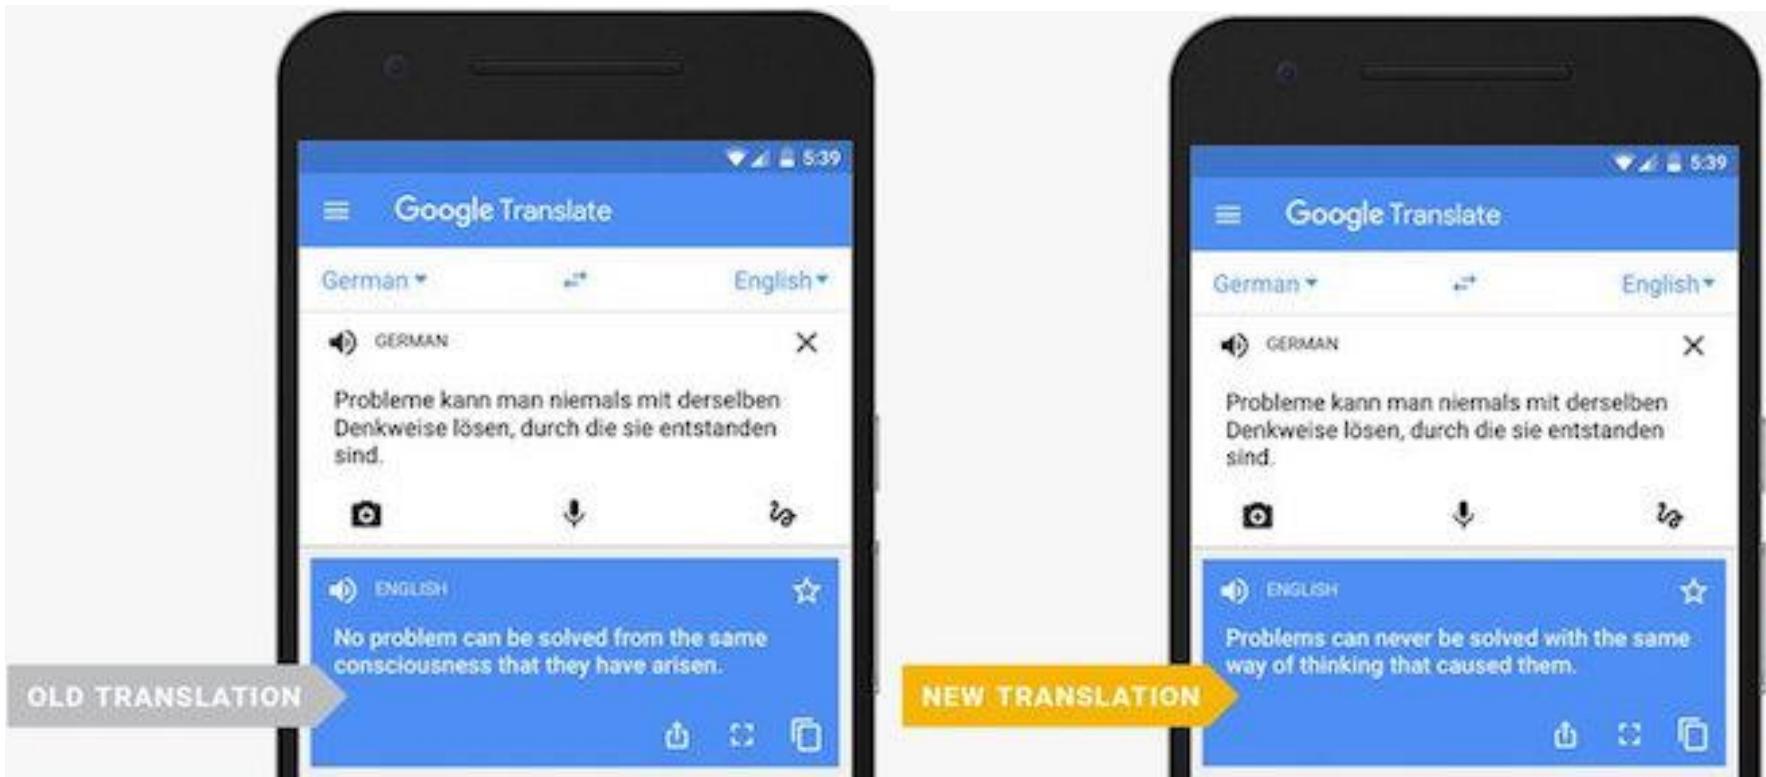
Verb tense



Country-Capital

<https://www.tensorflow.org/tutorials/word2vec>

Machine Translation



<https://www.pcmag.com/news/349610/google-expands-neural-networks-for-language-translation>

Text synthesis

Content: Two dogs play by a tree.

Style: **happily, love**



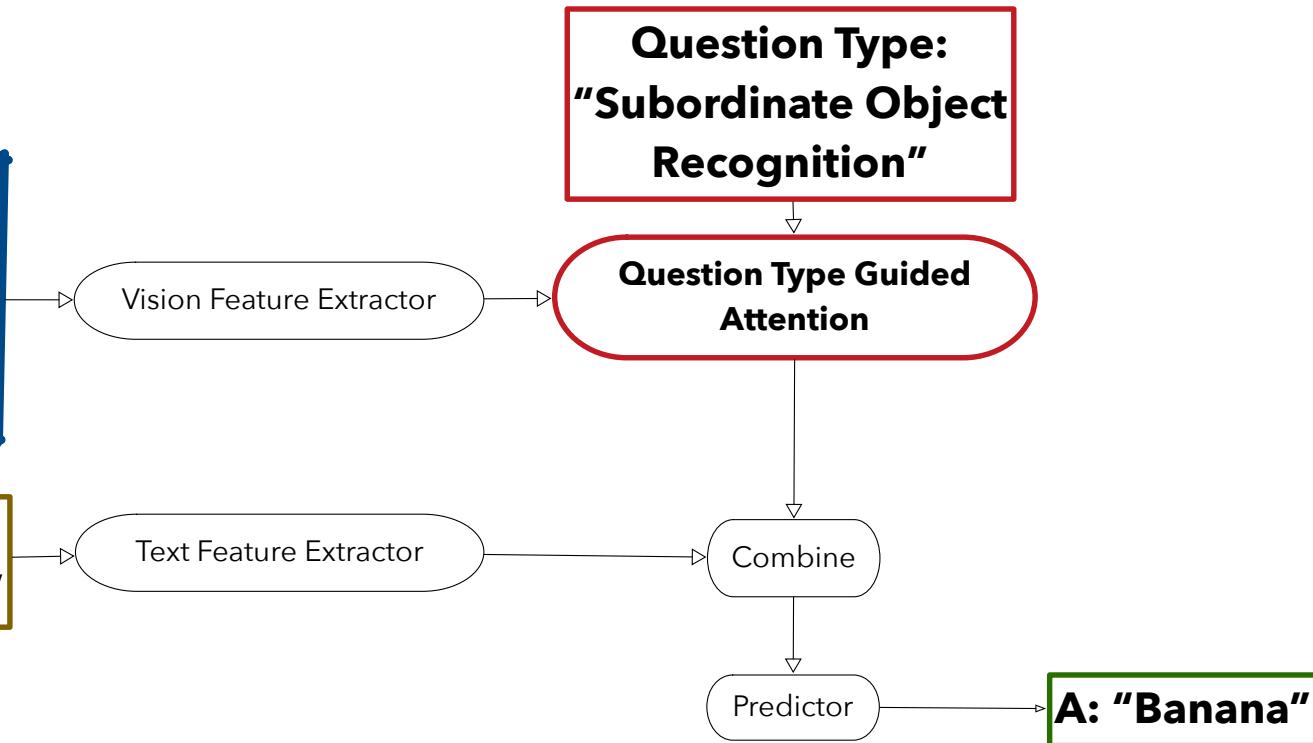
Two dogs **in love** play **happily** by a tree.

Li et al, NACCL, 2018

Question answering



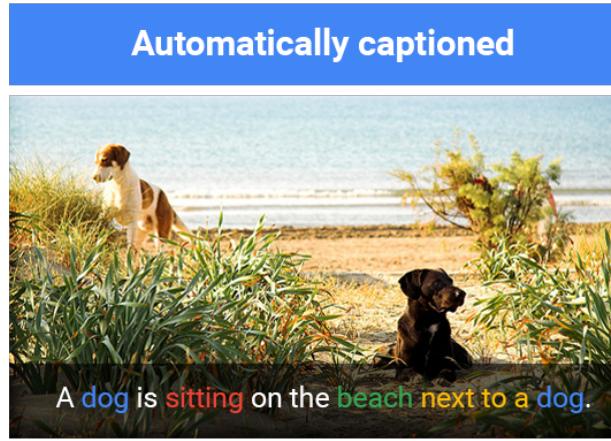
**Q: "What's her
mustache made of?"**



Shi et al, 2018, Arxiv

Image captioning

Human captions from the training set



Shallue et al, 2016

<https://ai.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>



Software

Tools

d2l.ai/chapter_crashcourse/install.html

- **Python**
 - Everyone is using it in machine learning & data science
 - Conda package manager (for simplicity)
- **Jupyter**
 - So much easier to keep track of your experiments
 - Obviously you should put longer code into modules
- **Reveal (for notebook slides)**
`conda install -c conda-forge rise`
- **Apache MXNet Gluon**
 - Scalability & ease of use
 - Imperative interface

Laptop / Desktop / Generic Cloud with Linux

- **Conda**

```
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh  
sh Miniconda3-latest-Linux-x86_64.sh  
mkdir d2l-en  
cd d2l-en  
curl https://www.d2l.ai/d2l-en.zip -o d2l-en.zip  
unzip d2l-en.zip  
rm d2l-en.zip
```

- **Install environment**

```
conda env create -f environment.yml  
cd d2l-en  
source activate gluon  
jupyter notebook
```

- **GPU**

- Install NVIDIA drivers / CUDA / CUDNN / TensorRT
- Update environment.yml to replace mxnet with mxnet-cu92

Laptop / Desktop with MacOS

- **Conda**

```
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-MacOS-x86_64.sh  
sh Miniconda3-latest-MacOS-x86_64.sh  
mkdir d2l-en  
cd d2l-en  
curl https://www.d2l.ai/d2l-en.zip -o d2l-en.zip  
unzip d2l-en.zip  
rm d2l-en.zip
```

- **Install environment**

```
conda env create -f environment.yml  
cd d2l-en  
source activate gluon  
jupyter notebook
```

- **GPU**

- Install NVIDIA drivers / CUDA / CUDNN / TensorRT
- Update environment.yml to replace mxnet with mxnet-cu92

Laptop / Desktop with Windows

- **Conda**

```
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Windows-x86_64.exe
Click on Miniconda3-latest-Windows-x86_64.exe
mkdir d2l-en
cd d2l-en
curl https://www.d2l.ai/d2l-en.zip -o d2l-en.zip
unzip d2l-en.zip
rm d2l-en.zip
```

- **Install environment**

```
conda env create -f environment.yml
cd d2l-en
activate gluon
jupyter notebook
```

- **GPU**

- Install NVIDIA drivers / CUDA / CUDNN / TensorRT
- Update environment.yml to replace mxnet with mxnet-cu92

AWS (via Deep Learning AMI)

- Install Base AMI (Ubuntu with NVIDIA drivers)
 - aws.amazon.com/machine-learning/amis/
 - aws.amazon.com/marketplace/pp/B077GFM7L7
- Follow steps for basic MXNet / Jupyter install as before
 - For P2 / P3 / G2 / G3 instances NVIDIA drivers default
 - Update the Conda environment
- SSH with port forwarding
 - d2l.ai/chapter_appendix/jupyter.html
 - d2l.ai/chapter_appendix/aws.html (setting up AWS account)



Colab

- Go to colab.research.google.com
- Activate the GPU supported runtime (this is a K80 GPU)
- Install MXNet

```
!pip install mxnet-cu92  
!pip install d2l
```



Linear Algebra

Scalars



- **Simple operations**

$$c = a + b$$

$$c = a \cdot b$$

$$c = \sin a$$

- **Length**

$$|a| = \begin{cases} a & \text{if } a > 0 \\ -a & \text{otherwise} \end{cases}$$

$$|a + b| \leq |a| + |b|$$

$$|a \cdot b| = |a| \cdot |b|$$

Vectors



- **Simple operations**

$$c = a + b \quad \text{where } c_i = a_i + b_i$$

$$c = \alpha \cdot b \quad \text{where } c_i = \alpha b_i$$

$$c = \sin a \quad \text{where } c_i = \sin a_i$$

- **Length**

Definition of a vector space

$$\|a\|_2 = \left[\sum_{i=1}^m a_i^2 \right]^{\frac{1}{2}}$$

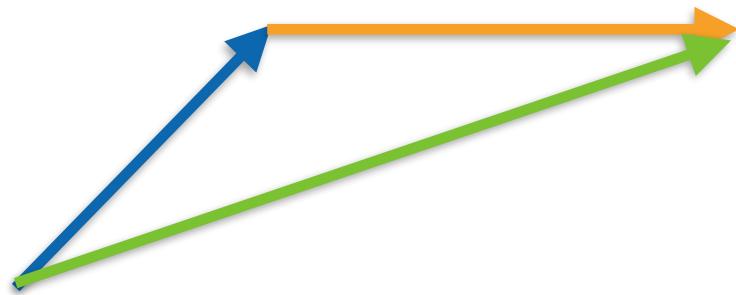
$$\|a\| \geq 0 \text{ for all } a$$

$$\|a + b\| \leq \|a\| + \|b\|$$

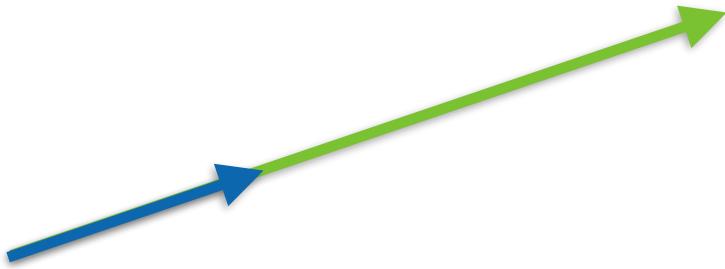
$$\|a \cdot b\| = |a| \cdot \|b\|$$

Definition of a norm

Vectors



$$c = a + b$$



$$c = \alpha \cdot b$$

Mathematician's 'parallel for all do'

Vectors

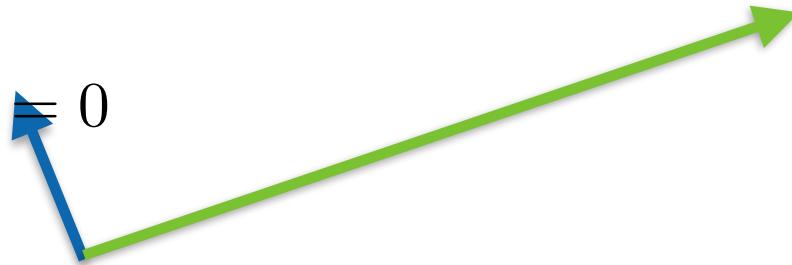


- Dot product

$$a^\top b = \sum_i a_i b_i$$

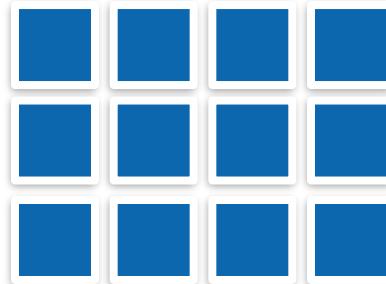
- Orthogonality

$$a^\top b = \sum_i a_i b_i = 0$$



(e.g. if we have two vectors that are orthogonal with a third, their linear combination is it, too)

Matrices



- **Simple operations**

$$C = A + B \quad \text{where } C_{ij} = A_{ij} + B_{ij}$$

$$C = \alpha \cdot B \quad \text{where } C_{ij} = \alpha B_{ij}$$

$$C = \sin A \quad \text{where } C_{ij} = \sin A_{ij}$$

- **Functional Analysis 101**

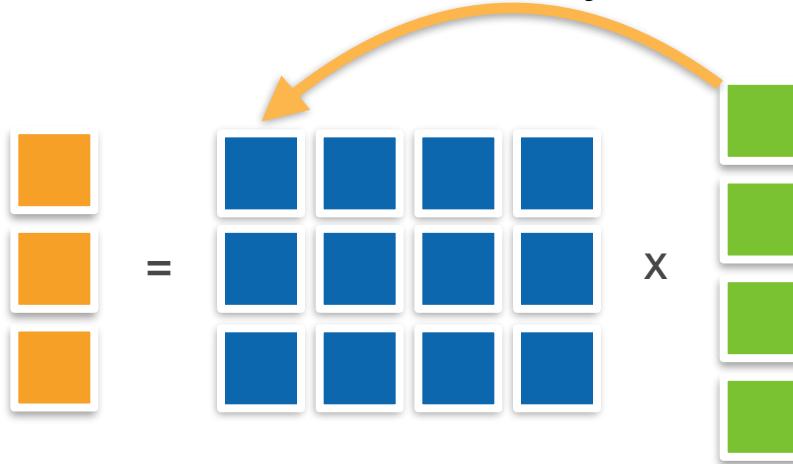
vector = function, matrix = linear operator

most theorems work sort-of in infinite dimensional spaces

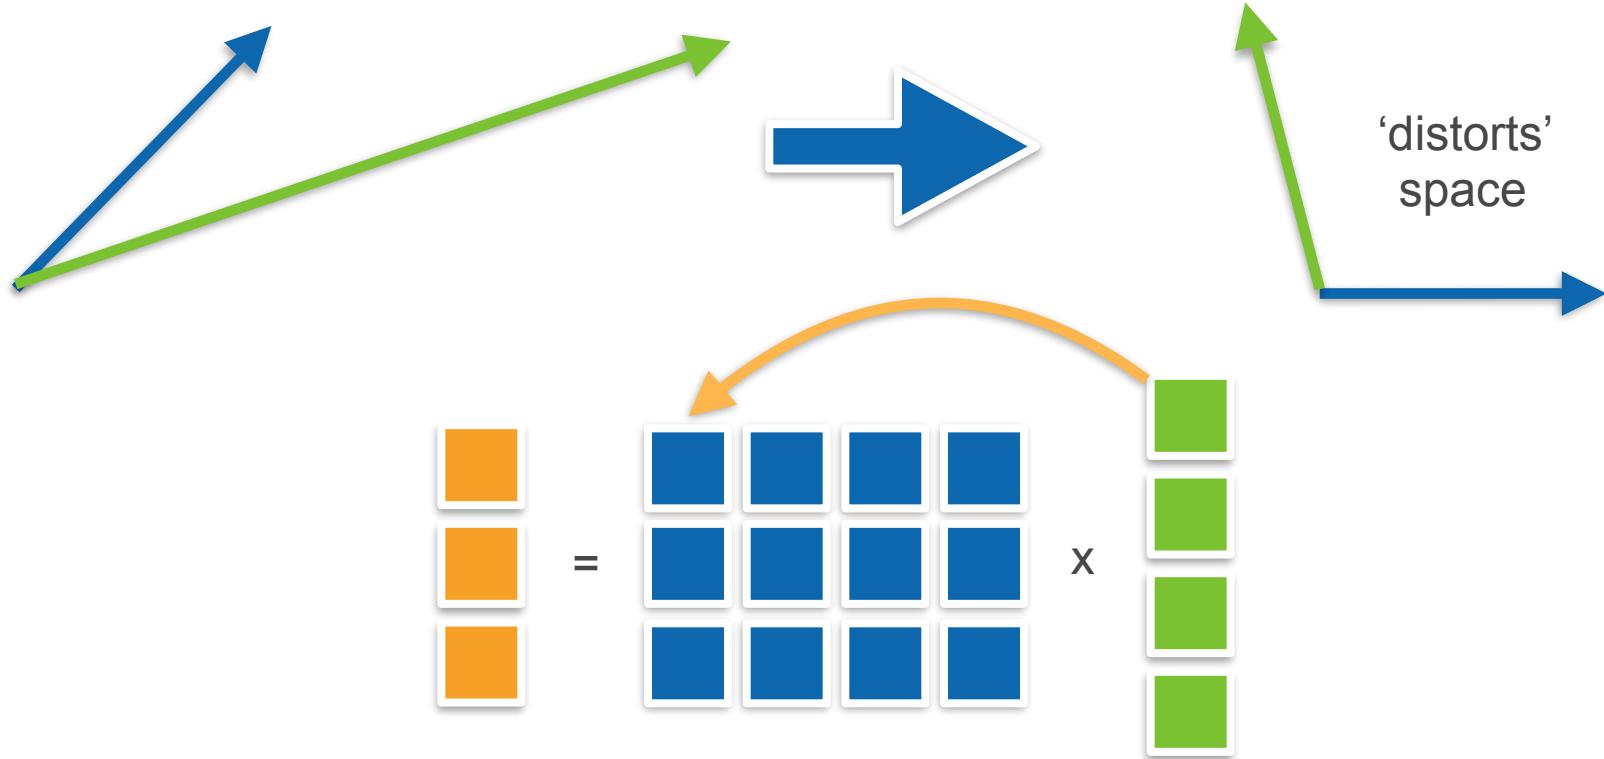
Matrices

- **Multiplications (matrix vector)**

$$c = Ab \text{ where } c_i = \sum_j A_{ij} b_j$$



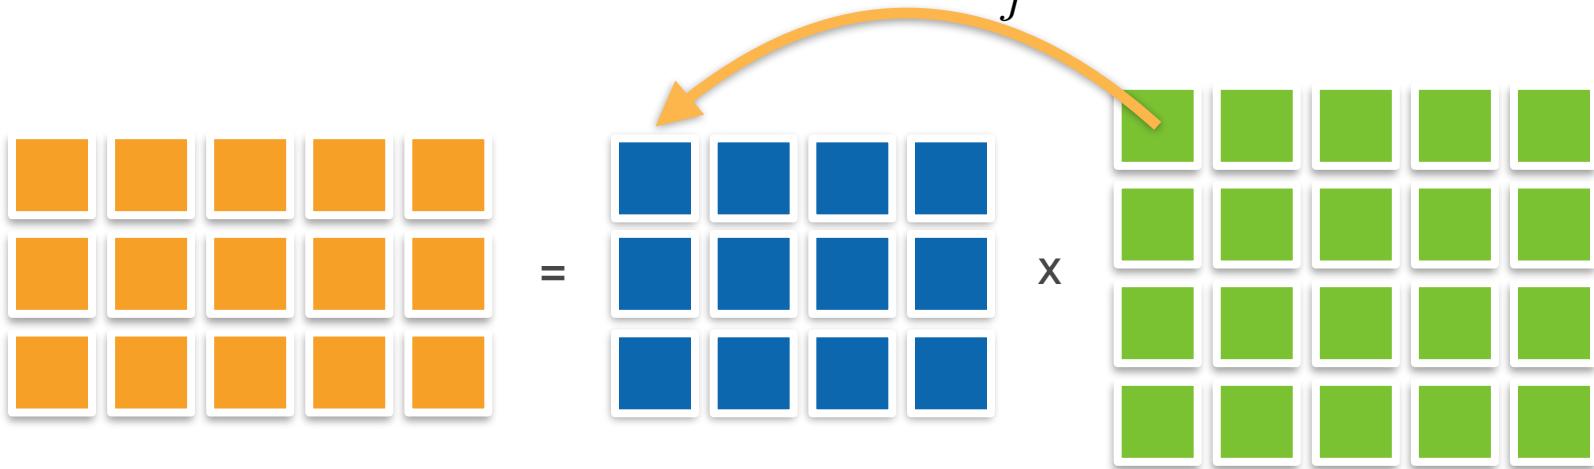
Matrices



Matrices

- **Multiplications (matrix matrix)**

$$C = AB \text{ where } C_{ik} = \sum_j A_{ij} B_{jk}$$



Matrices

- **Norms**

$$c = A \cdot b \text{ hence } \|c\| \leq \|A\| \cdot \|b\|$$

- Choices depending on how to measure length of b and c

- **Popular norms**

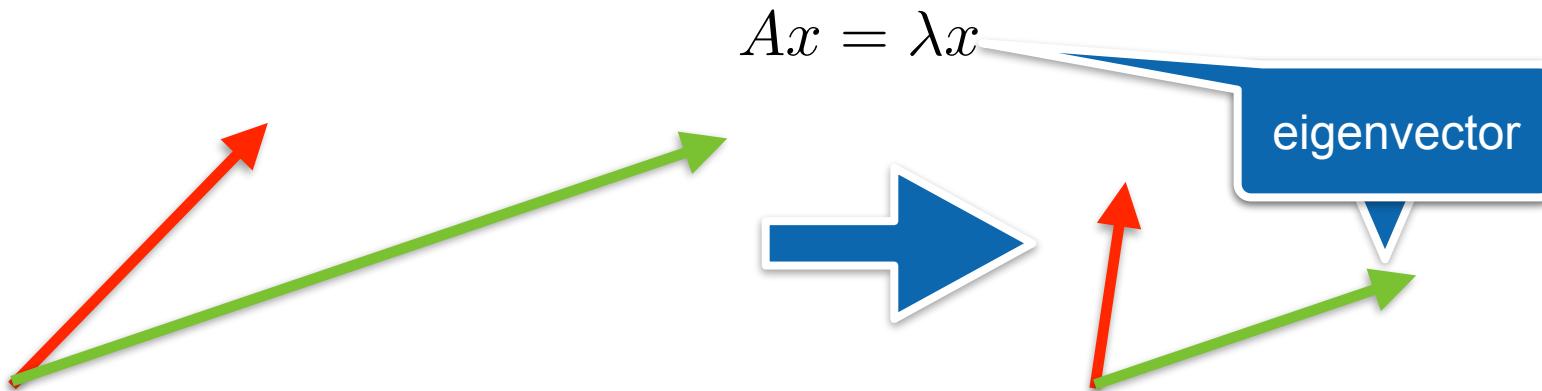
- Frobenius norm

$$\|A\|_{\text{Frob}} = \left[\sum_{ij} A_{ij}^2 \right]^{\frac{1}{2}}$$

- H-infinity norm (or often just matrix norm) as smallest value for which condition holds for Euclidean norms on vectors.

Matrices

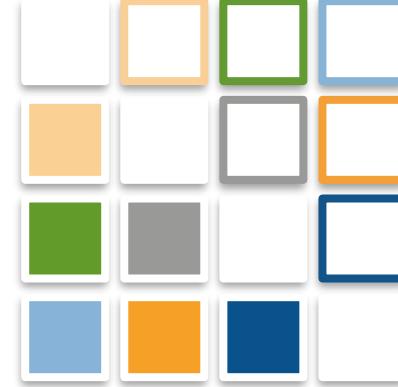
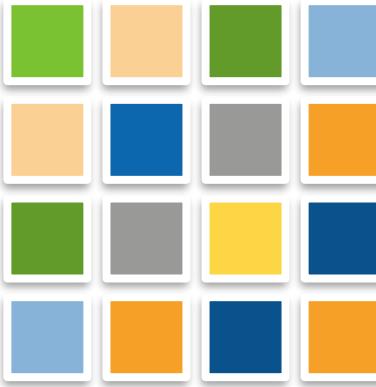
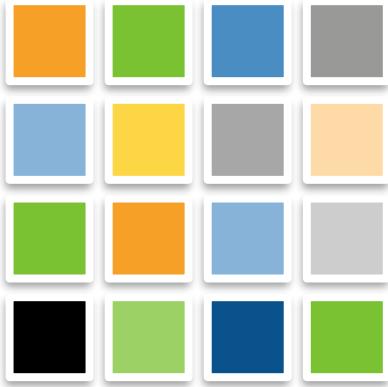
- **Eigenvectors and eigenvalue**
 - Vectors that aren't changed by the matrix



- For symmetric matrices we can always find this

Special Matrices

- **Symmetric, antisymmetric** $A_{ij} = A_{ji}$ and $A_{ij} = -A_{ji}$



- **Positive definite**

$$\|x\|^2 = x^\top x \geq 0 \text{ generalizes to } x^\top Ax \geq 0$$

(all positive eigenvalues)

Special Matrices

- **Orthogonal Matrices**

- All rows of the matrix are orthogonal to each other
- All rows of the matrix have unit length

$$U \text{ with } \sum_j U_{ij} U_{kj} = \delta_{ik}$$

- Rewrite in matrix form

$$UU^\top = \mathbf{1}$$

Show that
 $U^\top U = \mathbf{1}$

- **Permutation Matrices**

$$P \text{ where } P_{ij} = 1 \text{ if and only if } j = \pi(i)$$

Show that P is
orthogonal



GPUs love matrices and vectors
(they have many processors)



ndarray

N-dimensional Array Examples

- N-dimensional array, short for ndarray, is the main data structure for machine learning and neural networks

0-d (scalar)



1.0

A class label

1-d (vector)



[1.0, 2.7, 3.4]

A feature vector

2-d (matrix)

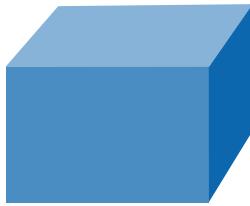


[[1.0, 2.7, 3.4],
 [5.0, 0.2, 4.6],
 [4.3, 8.5, 0.2]]

A example-by-feature matrix

ND Array Examples, cont

3-d



```
[[[0.1, 2.7, 3.4]  
 [5.0, 0.2, 4.6]  
 [4.3, 8.5, 0.2]]  
 [[3.2, 5.7, 3.4]  
 [5.4, 6.2, 3.2]  
 [4.1, 3.5, 6.2]]]
```

A RGB image
(width x height
x channels)

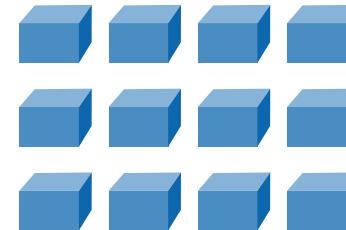
4-d



```
[[[[. . .  
 . . .  
 . . .]]]
```

A batch of
RGB images
(batch-size x
width x height
x channels)

5-d



```
[[[[. . .  
 . . .  
 . . .]]]
```

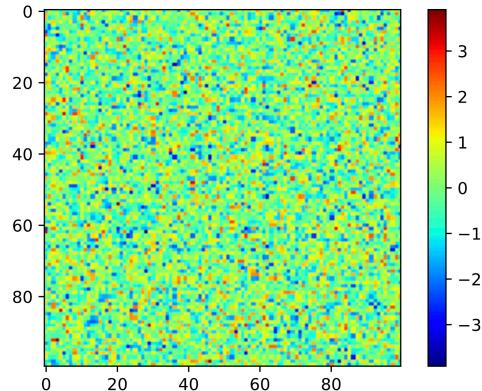
A batch of videos
(batch-size x time x
width x height x
channels)

Create Arrays

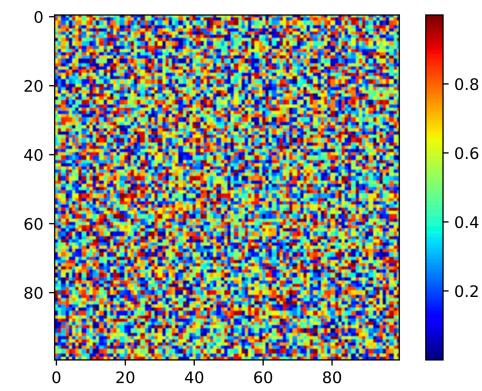
- Create arrays with
 - A shape, e.g. 3-by-4 matrix
 - Data type for each element, e.g. float
 - Element values, e.g all 0s, or random values

100-by-100 matrix
with elements
generated from

A normal distribution



A uniform distribution



Access Elements

An element: [1, 2]

	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12
3	13	14	15	16

A row: [1, :]

	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12
3	13	14	15	16

A column: [1, :]

	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12
3	13	14	15	16

	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12
3	13	14	15	16

	0	1	2	3
0	1	2	3	4
1	5	6	7	8
2	9	10	11	12
3	13	14	15	16

Pointers

- d2l.ai/chapter_crashcourse/ndarray.html
- beta.mxnet.io/guide/crash-course/1-ndarray.html
- Gilbert Strang's Linear Algebra course
github.com/juanklopper/MIT_OCW_Linear_Algebra_18_06
- Berkeley HPC course (GPU sections)
sites.google.com/lbl.gov/cs267-spr2018/

Summary

- **Logistics**
Homework, Project, Exam, Hours, User group
- **Deep Learning**
Quick overview
- **Software**
Installation, cloud
- **Linear Algebra**
Basic notation
- **NDArray**